

```
In [115]: import pandas as pd # For handling data
import numpy as np # For numerical computations
import matplotlib.pyplot as plt # For visualizations
import math
import seaborn as sns # For better visualizations
from sklearn.model_selection import train_test_split # To split data
from sklearn.preprocessing import OneHotEncoder, StandardScaler # For preprocessing
from sklearn.ensemble import RandomForestRegressor # For building model
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score # For evaluation
import warnings
with warnings.catch_warnings(): warnings.simplefilter('ignore')
warnings.filterwarnings('ignore')
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from pptx import Presentation
from pptx.util import Inches
from sklearn.ensemble import GradientBoostingRegressor
```

Step 1: Load the Dataset

```
In [40]: mobile_data=pd.read_csv('Mobile_data.csv')
```

```
In [42]: mobile_data
```

Out[42]:

	Unnamed: 0	Model	Colour	Memory	RAM	Battery_	Rear Camera	Front Camera	AI Lens	Mobile Height	Processor_	Prize
0	0	Infinix SMART 7	Night Black	64	4	6000	13MP	5MP	1	16.76	Unisoc Spreadtrum SC9863A1	7,299
1	1	Infinix SMART 7	Azure Blue	64	4	6000	13MP	5MP	1	16.76	Unisoc Spreadtrum SC9863A1	7,299
2	2	MOTOROLA G32	Mineral Gray	128	8	5000	50MP	16MP	0	16.64	Qualcomm Snapdragon 680	11,999
3	3	POCO C50	Royal Blue	32	2	5000	8MP	5MP	0	16.56	Mediatek Helio A22	5,649
4	4	Infinix HOT 30i	Marigold	128	8	5000	50MP	5MP	1	16.76	G37	8,999
...
536	637	SAMSUNG Galaxy S23 5G	Cream	256	8	3900	50MP	12MP	0	15.49	Qualcomm Snapdragon 8 Gen 2	79,999
537	638	LAVA Z21	Cyan	32	2	3100	5MP	2MP	0	12.70	Octa Core	5,998
538	639	Tecno Spark 8T	Turquoise Cyan	64	4	5000	50MP	8MP	0	16.76	MediaTek Helio G35	9,990
539	641	SAMSUNG Galaxy A54 5G	Awesome Lime	128	8	5000	50MP	32MP	0	16.26	Exynos 1380, Octa Core	38,999
540	642	OPPO A77	Sky Blue	128	4	5000	50MP	8MP	0	16.66	Mediatek Helio G35	15,999

541 rows × 12 columns

In [9]: `print(mobile_data.info())`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541 entries, 0 to 540
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Unnamed: 0      541 non-null   int64
 1   Model           541 non-null   object
 2   Colour         541 non-null   object
 3   Memory         541 non-null   int64
 4   RAM            541 non-null   int64
 5   Battery_       541 non-null   int64
 6   Rear Camera    541 non-null   object
 7   Front Camera   541 non-null   object
 8   AI Lens        541 non-null   int64
 9   Mobile Height  541 non-null   float64
10  Processor_     541 non-null   object
11  Prize          541 non-null   object
dtypes: float64(1), int64(5), object(6)
memory usage: 50.8+ KB
None
```

```
In [11]: mobile_data.head()
```

Out[11]:

	Unnamed: 0	Model	Colour	Memory	RAM	Battery_	Rear Camera	Front Camera	AI Lens	Mobile Height	Processor_	Prize
0	0	Infinix SMART 7	Night Black	64	4	6000	13MP	5MP	1	16.76	Unisoc Spreadtrum SC9863A1	7,299
1	1	Infinix SMART 7	Azure Blue	64	4	6000	13MP	5MP	1	16.76	Unisoc Spreadtrum SC9863A1	7,299
2	2	MOTOROLA G32	Mineral Gray	128	8	5000	50MP	16MP	0	16.64	Qualcomm Snapdragon 680	11,999
3	3	POCO C50	Royal Blue	32	2	5000	8MP	5MP	0	16.56	Mediatek Helio A22	5,649
4	4	Infinix HOT 30i	Marigold	128	8	5000	50MP	5MP	1	16.76	G37	8,999

In [13]: `mobile_data.tail()`

Out[13]:

	Unnamed: 0	Model	Colour	Memory	RAM	Battery_	Rear Camera	Front Camera	AI Lens	Mobile Height	Processor_	Prize
536	637	SAMSUNG Galaxy S23 5G	Cream	256	8	3900	50MP	12MP	0	15.49	Qualcomm Snapdragon 8 Gen 2	79,999
537	638	LAVA Z21	Cyan	32	2	3100	5MP	2MP	0	12.70	Octa Core	5,998
538	639	Tecno Spark 8T	Turquoise Cyan	64	4	5000	50MP	8MP	0	16.76	MediaTek Helio G35	9,990
539	641	SAMSUNG Galaxy A54 5G	Awesome Lime	128	8	5000	50MP	32MP	0	16.26	Exynos 1380, Octa Core	38,999
540	642	OPPO A77	Sky Blue	128	4	5000	50MP	8MP	0	16.66	Mediatek Helio G35	15,999

```
In [15]: mobile_data.shape
```

```
Out[15]: (541, 12)
```

```
In [9]: # Check for missing values  
missing_values = mobile_data.isnull().sum()  
missing_values[missing_values > 0] # Show only columns with missing values
```

```
Out[9]: Series([], dtype: int64)
```

```
In [44]: mobile_data.isnull().sum()
```

```
Out[44]: Unnamed: 0      0  
Model      0  
Colour     0  
Memory     0  
RAM        0  
Battery_   0  
Rear Camera 0  
Front Camera 0  
AI Lens    0  
Mobile Height 0  
Processor_ 0  
Prize      0  
dtype: int64
```

```
In [46]: mobile_data.drop(columns='Unnamed: 0', inplace=True, errors='ignore')  
mobile_data
```

Out[46]:

	Model	Colour	Memory	RAM	Battery_	Rear Camera	Front Camera	AI Lens	Mobile Height	Processor_	Prize
0	Infinix SMART 7	Night Black	64	4	6000	13MP	5MP	1	16.76	Unisoc Spreadtrum SC9863A1	7,299
1	Infinix SMART 7	Azure Blue	64	4	6000	13MP	5MP	1	16.76	Unisoc Spreadtrum SC9863A1	7,299
2	MOTOROLA G32	Mineral Gray	128	8	5000	50MP	16MP	0	16.64	Qualcomm Snapdragon 680	11,999
3	POCO C50	Royal Blue	32	2	5000	8MP	5MP	0	16.56	Mediatek Helio A22	5,649
4	Infinix HOT 30i	Marigold	128	8	5000	50MP	5MP	1	16.76	G37	8,999
...
536	SAMSUNG Galaxy S23 5G	Cream	256	8	3900	50MP	12MP	0	15.49	Qualcomm Snapdragon 8 Gen 2	79,999
537	LAVA Z21	Cyan	32	2	3100	5MP	2MP	0	12.70	Octa Core	5,998
538	Tecno Spark 8T	Turquoise Cyan	64	4	5000	50MP	8MP	0	16.76	MediaTek Helio G35	9,990
539	SAMSUNG Galaxy A54 5G	Awesome Lime	128	8	5000	50MP	32MP	0	16.26	Exynos 1380, Octa Core	38,999
540	OPPO A77	Sky Blue	128	4	5000	50MP	8MP	0	16.66	Mediatek Helio G35	15,999

541 rows × 11 columns

```
In [48]: # Check for duplicate rows
print(f"Duplicate rows: {mobile_data.duplicated().sum()}")
```

```
# Remove duplicates
mobile_data= mobile_data.drop_duplicates()
```

Duplicate rows: 10

```
In [50]: # Check for duplicate rows after the duplicate are dropped
print(f"Duplicate rows: {mobile_data.duplicated().sum()}")
```

Duplicate rows: 0

```
In [38]: #Fill categorical missing values using the most frequent value:
# Fill missing values based on column types
for column in mobile_data.columns:
    if mobile_data[column].dtype == "object": # If categorical
        mobile_data[column].fillna(mobile_data[column].mode()[0], inplace=True)
    else: # If numerical
        mobile_data[column].fillna(mobile_data[column].median(), inplace=True)

# Verify again
print(mobile_data.isnull().sum().sum()) # Should return 0 if all missing values are handled
```

0

```
In [52]: import matplotlib.pyplot as plt
import seaborn as sns
import math

# Identify numerical columns
numerical_cols = mobile_data.select_dtypes(include=["int64", "float64"]).columns

# Number of numerical columns
num_cols = len(numerical_cols)
cols = 2 # Two columns: one for histogram, one for boxplot
rows = math.ceil(num_cols) # Ensure enough rows

# Set a modern Seaborn theme
sns.set_style("darkgrid")

plt.figure(figsize=(14, rows * 5))
plt.suptitle("📊 Exploratory Data Analysis: Numerical Features", fontsize=16, fontweight="bold", color="darkred")

# Define color palette
hist_colors = ["royalblue", "seagreen", "purple", "darkorange", "crimson"]
```

```

box_colors = ["skyblue", "lightgreen", "mediumpurple", "gold", "lightcoral"]

# Loop through all numerical columns
for i, col in enumerate(numerical_cols):
    color_hist = hist_colors[i % len(hist_colors)]
    color_box = box_colors[i % len(box_colors)]


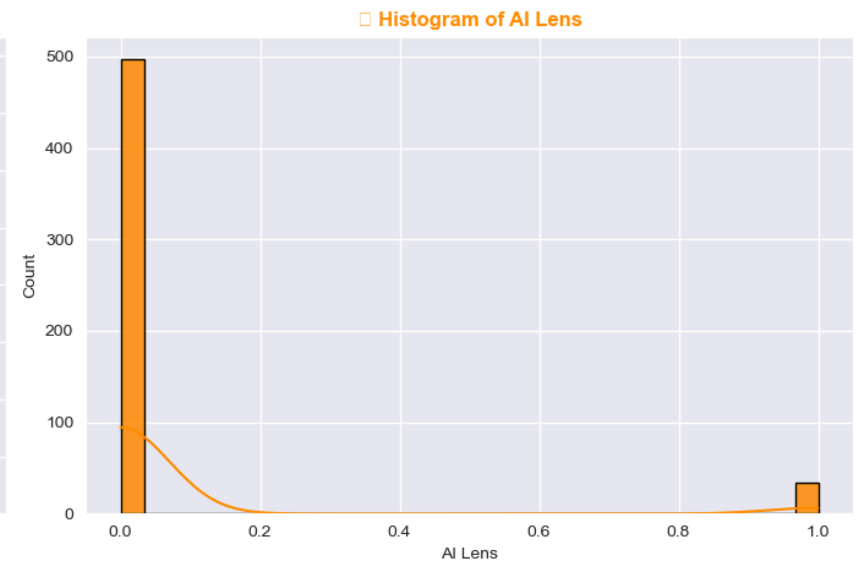
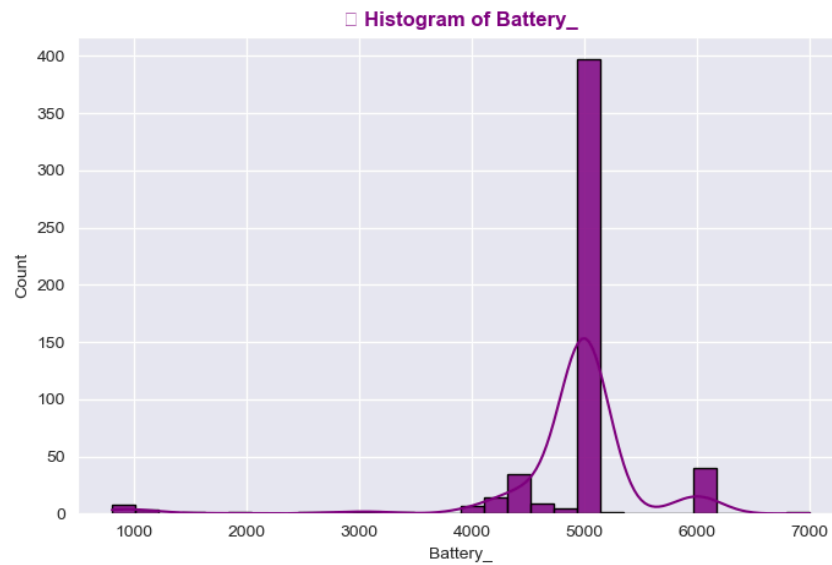
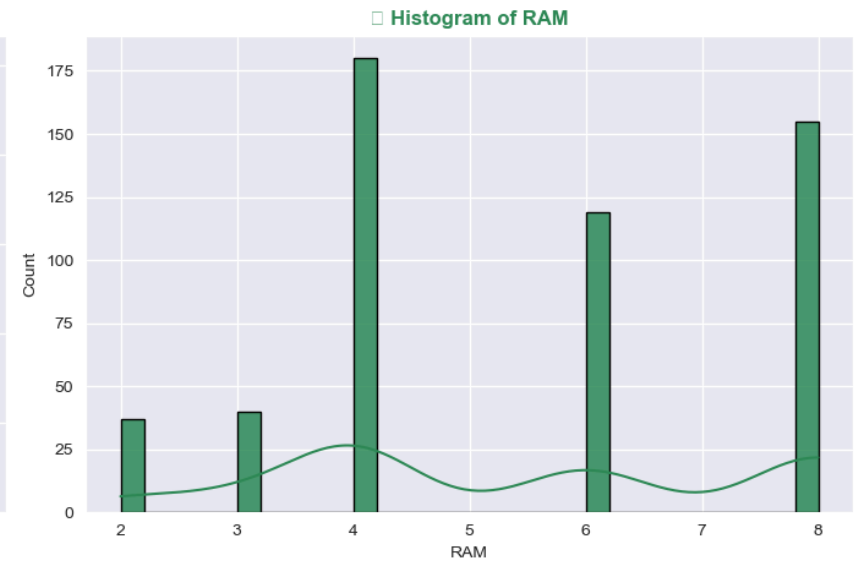

    # Create histogram
    plt.subplot(rows, cols, i + 1) # Adjusted subplot index
    sns.histplot(mobile_data[col], bins=30, kde=True, color=color_hist, edgecolor="black", alpha=0.85)
    plt.title(f"📊 Histogram of {col}", fontsize=12, fontweight="bold", color=color_hist)
    plt.xlabel(col, fontsize=10)
    plt.ylabel("Count", fontsize=10)


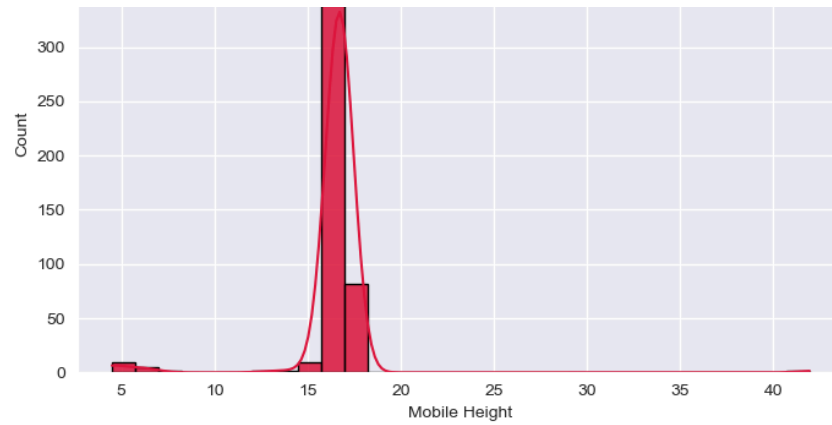
    # Create boxplot
    plt.subplot(rows, cols, i + num_cols + 1) # Adjusted subplot index
    sns.boxplot(x=mobile_data[col], color=color_box)
    plt.title(f"📦 Boxplot of {col}", fontsize=12, fontweight="bold", color=color_box)

plt.tight_layout(rect=[0, 0, 1, 0.95]) # Adjust layout to fit main title
plt.show()


```


Exploratory Data Analysis: Numerical Features






Boxplot of RAM



Boxplot of Battery_





Boxplot of AI Lens



Boxplot of Mobile Height





```
In [54]: import matplotlib.pyplot as plt
import seaborn as sns

# Identify categorical columns
categorical_cols = mobile_data.select_dtypes(include=["object"]).columns

# Set Seaborn style
sns.set_style("darkgrid")

# Generate count plots for categorical columns
for col in categorical_cols:
    plt.figure(figsize=(14, 6)) # Increase figure size

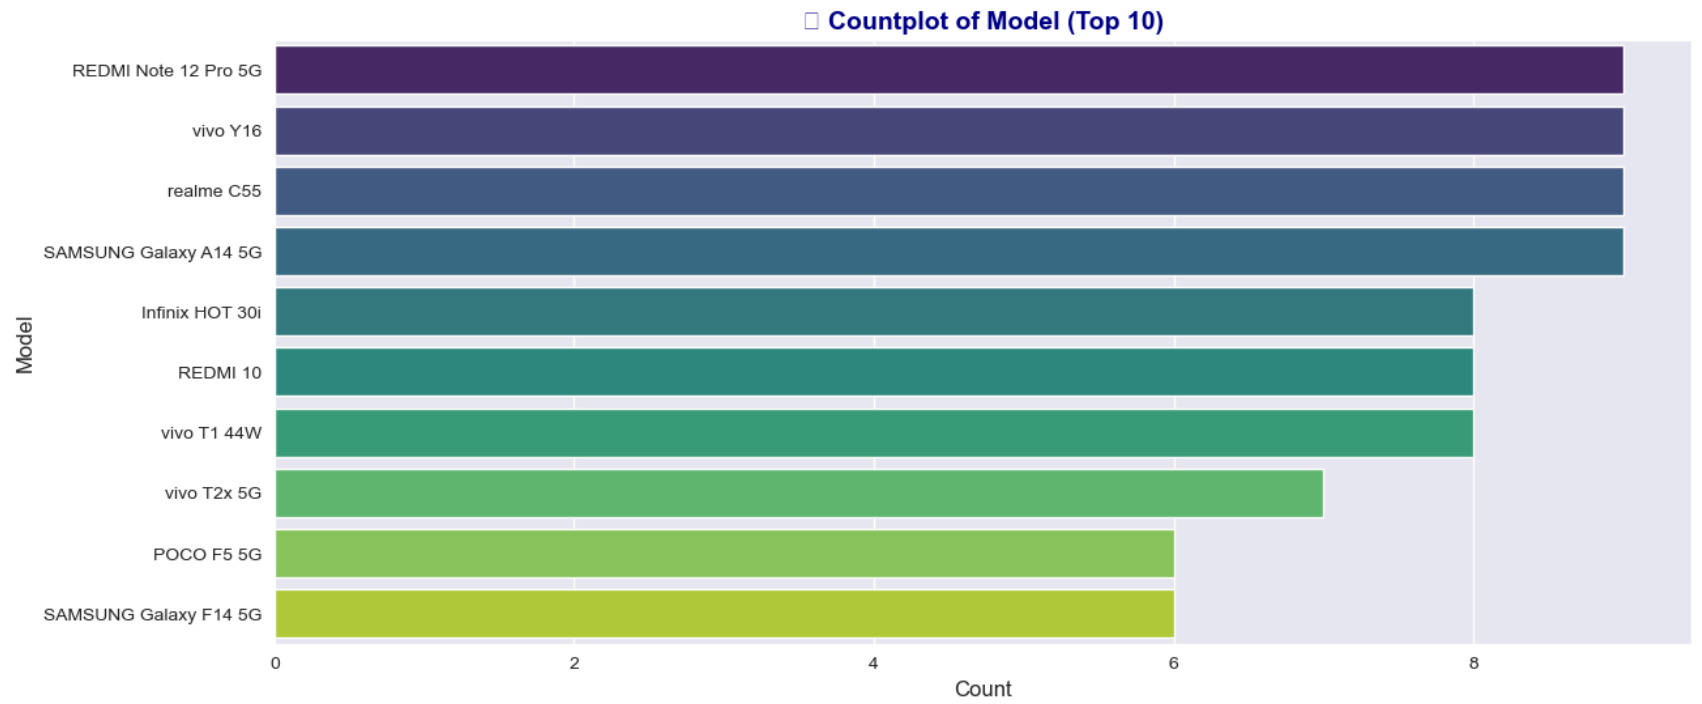
    # Show only top 10 categories to avoid clutter
    top_10_values = mobile_data[col].value_counts().nlargest(10).index
    filtered_data = mobile_data[mobile_data[col].isin(top_10_values)]

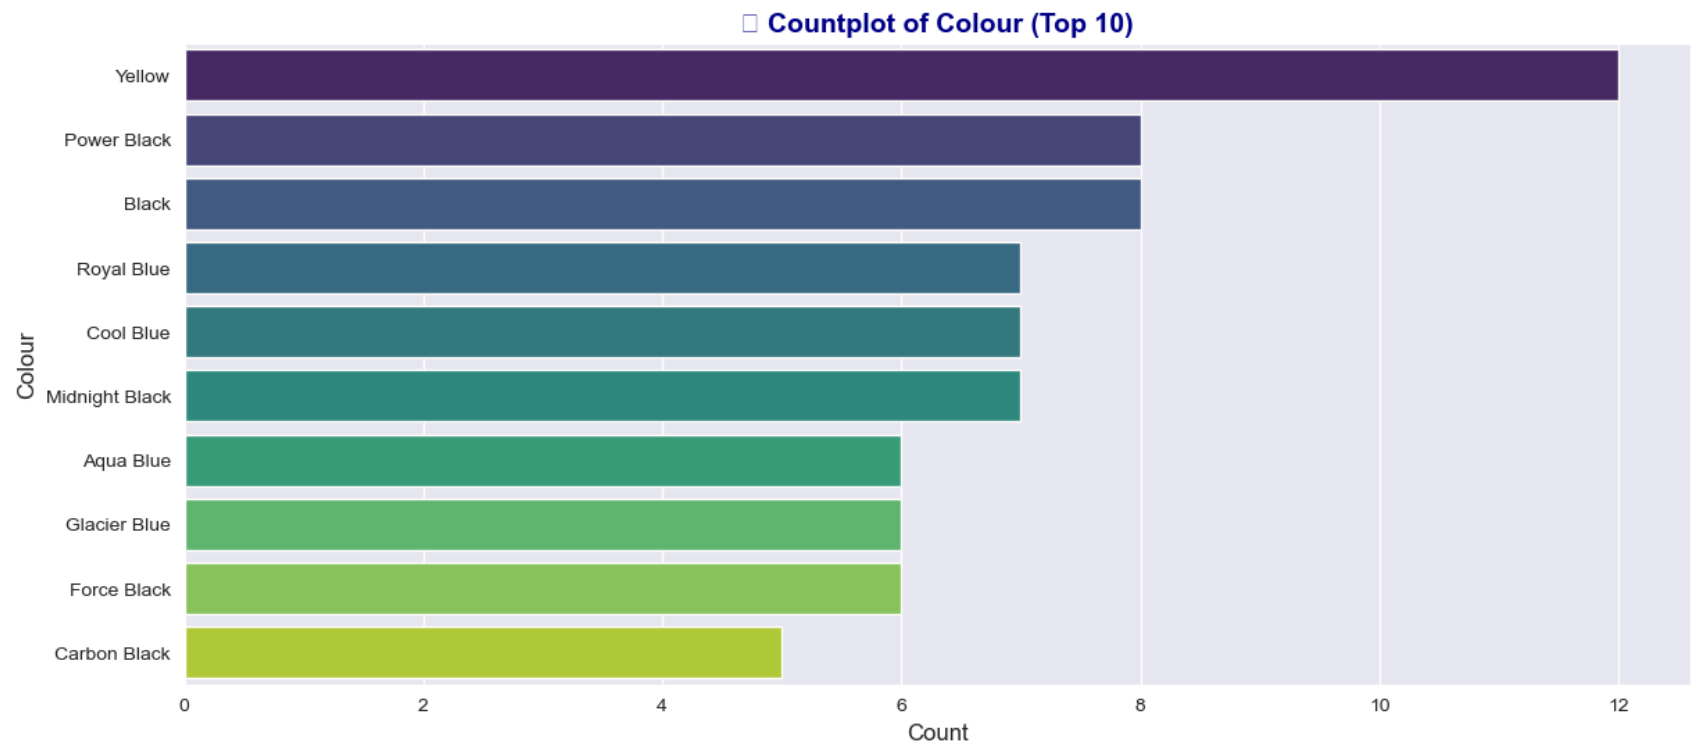
    sns.countplot(y=filtered_data[col], order=top_10_values, palette="viridis")

    plt.title(f"📊 Countplot of {col} (Top 10)", fontsize=14, fontweight="bold", color="darkblue")
    plt.ylabel(col, fontsize=12)
    plt.xlabel("Count", fontsize=12)


    plt.xticks(rotation=0) # Keep x labels horizontal
    plt.yticks(fontsize=10) # Increase y-axis font size


    plt.show()
```

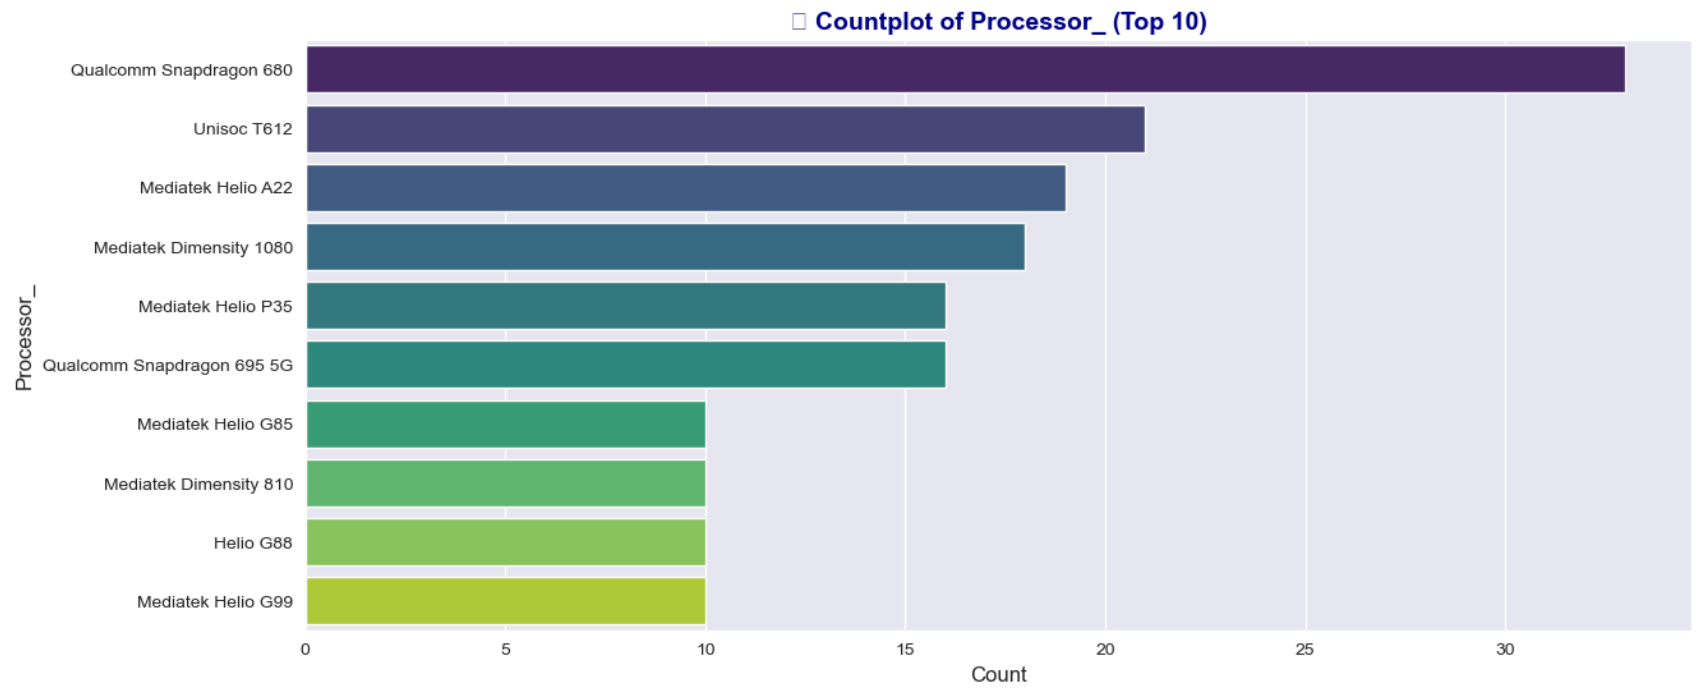





Countplot of Rear Camera (Top 10)












```
In [17]: # Identify numerical columns
numerical_cols = mobile_data.select_dtypes(include=["int64", "float64"]).columns

# Plot KDE for all numerical columns
for col in numerical_cols:
    plt.figure(figsize=(8, 5))
    sns.kdeplot(mobile_data[col], fill=True, color="blue")
    plt.title(f"Kernel Density Plot of {col}")
    plt.xlabel(col)
    plt.ylabel("Density")
    plt.show()
```


Kernel Density Plot of Memory

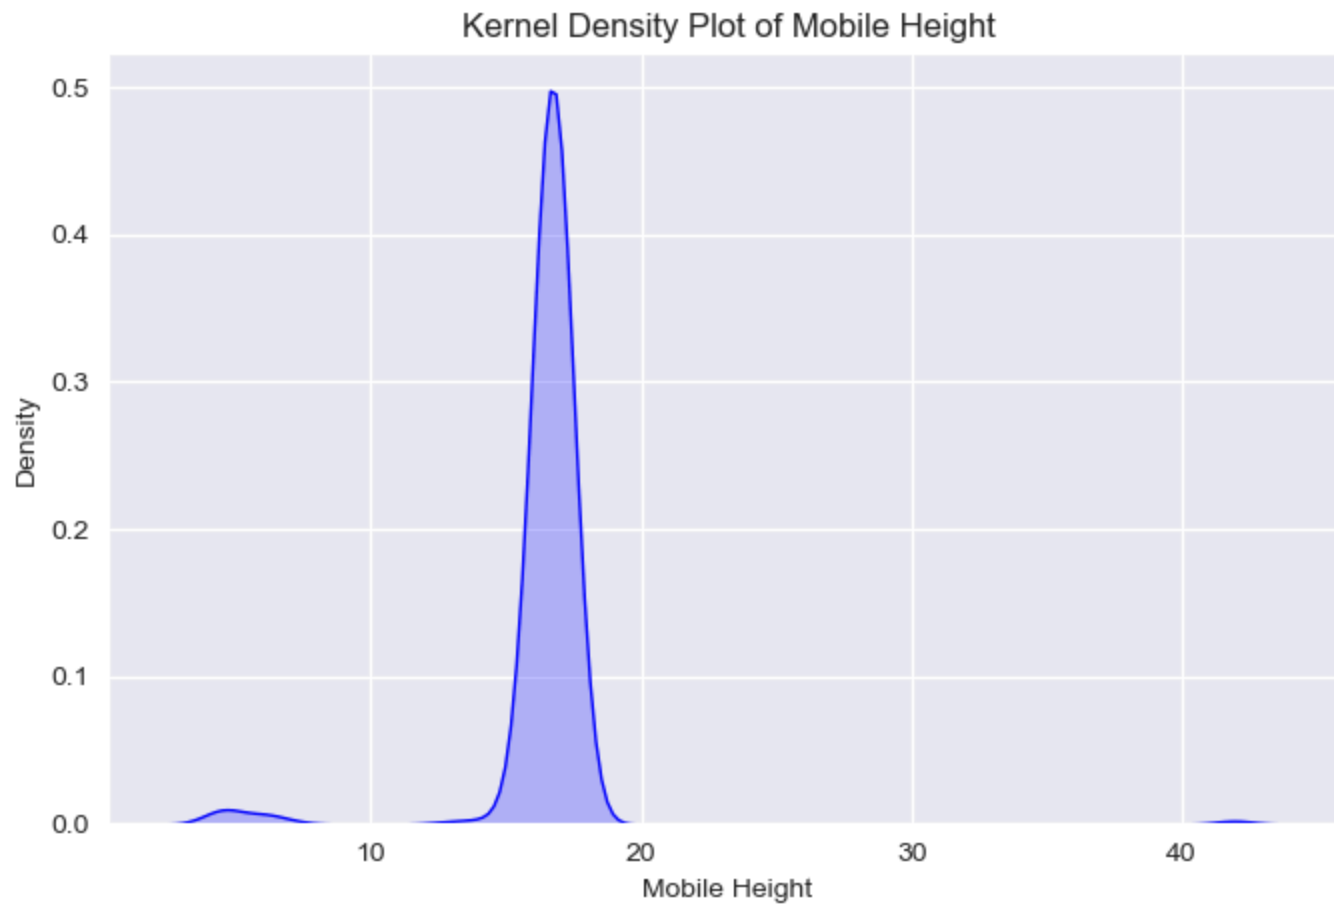






Kernel Density Plot of AI Lens





Pairplot (Scatterplot Matrix) with Categories


```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Set Seaborn style
sns.set_style("darkgrid")

# Select only numerical columns for correlation analysis
numerical_cols = mobile_data.select_dtypes(include=["int64", "float64"])
```

```
# 🔥 1. Correlation Matrix Heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(numerical_cols.corr(), annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
plt.title("🔍 Correlation Matrix Heatmap", fontsize=14, fontweight="bold", color="darkred")
plt.show()

# 🔥 2. Pairplot (Scatterplot Matrix)
pairplot_fig = sns.pairplot(numerical_cols.iloc[:, :5], diag_kind="kde", palette="husl")
pairplot_fig.fig.suptitle("📊 Scatterplot Matrix (First 5 Numerical Features)",
                          fontsize=14, fontweight="bold", color="darkblue", y=1.05) # Adjust title position
plt.show()
```



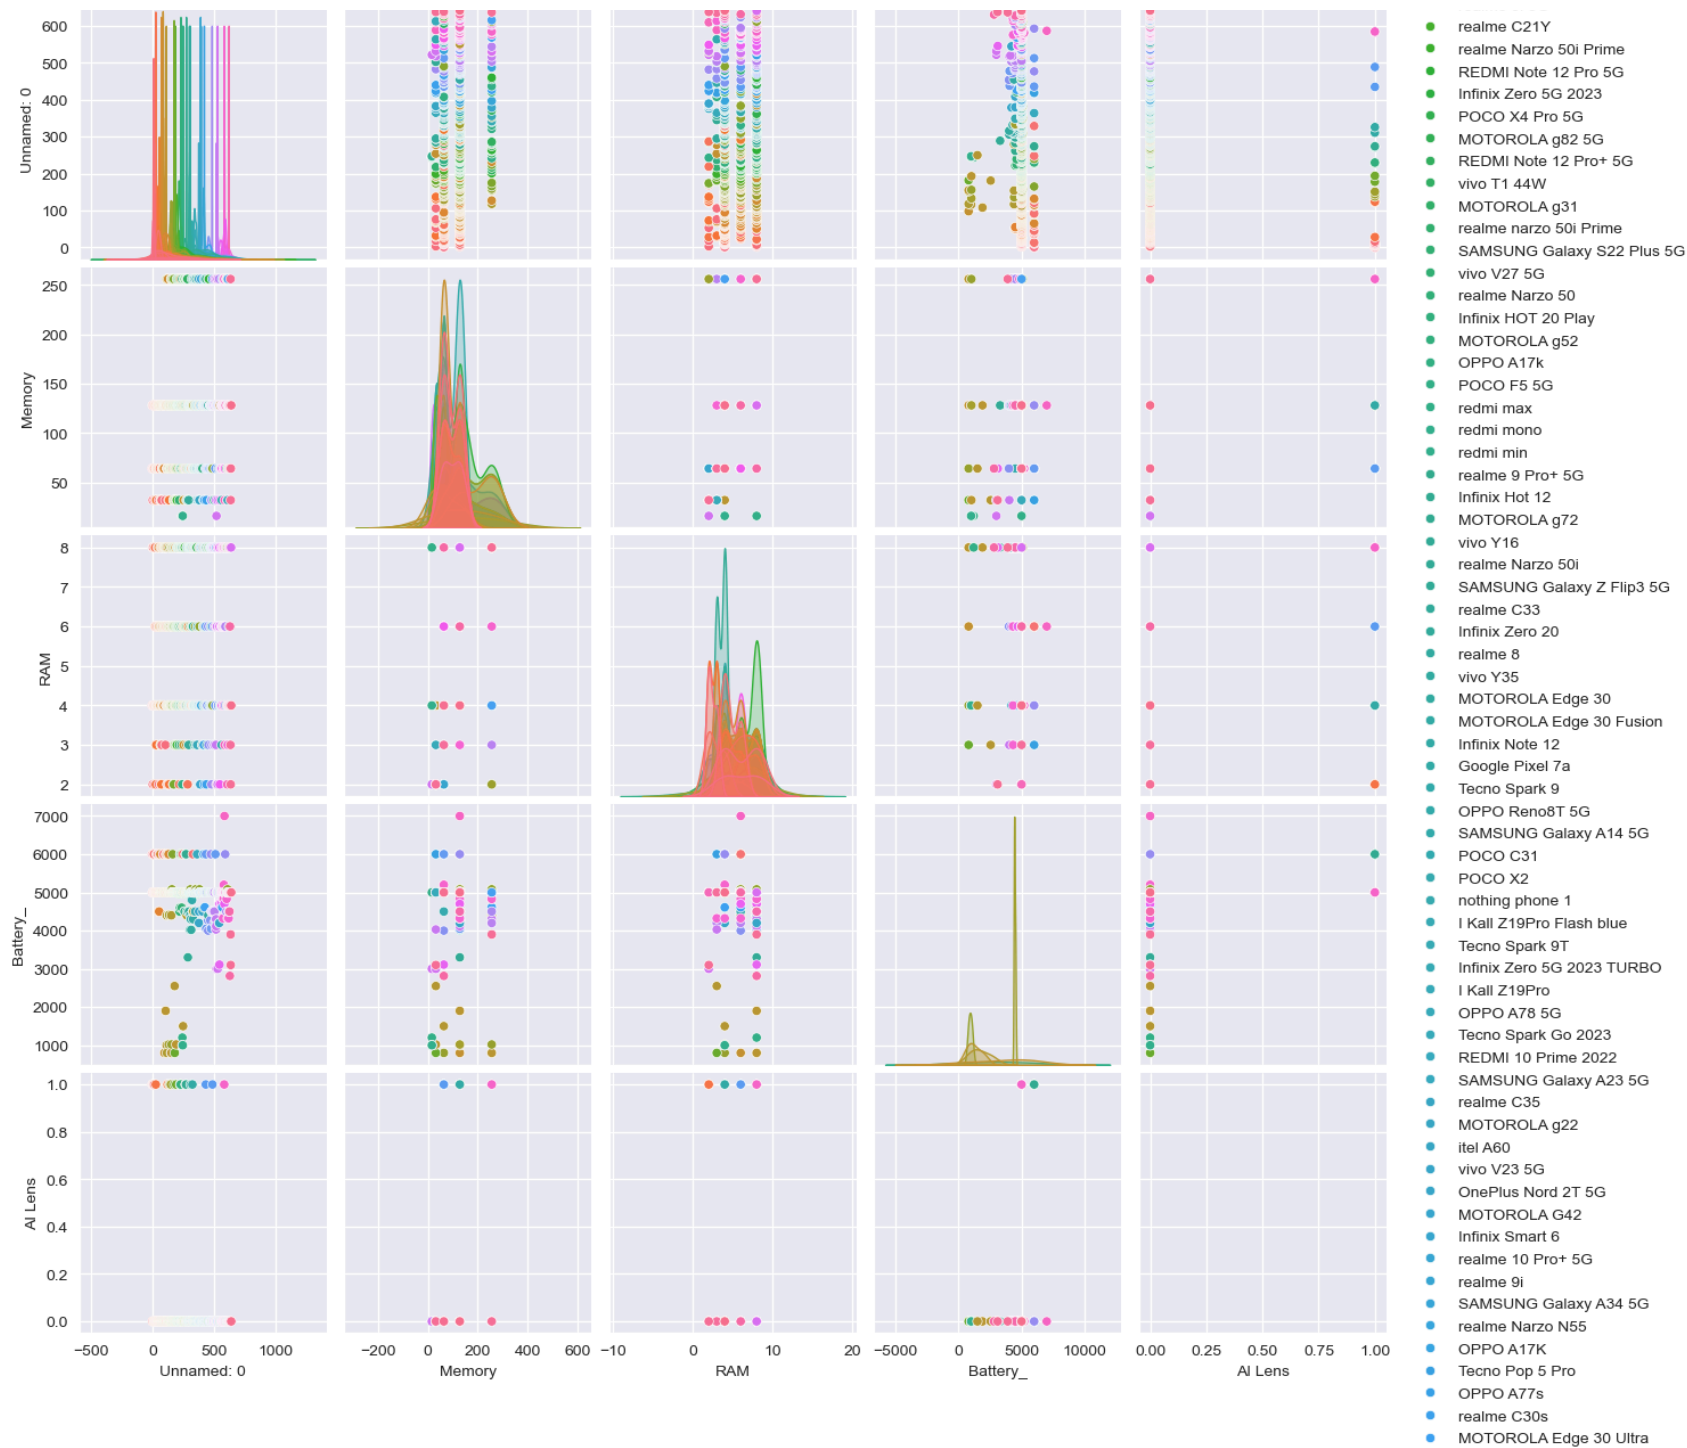
Correlation Heatmap with Categories


```
In [20]: # Select numerical features for analysis
numerical_cols = mobile_data.select_dtypes(include=["int64", "float64"]).columns

# Create a scatterplot matrix for selected features with category hue
pairplot_fig = sns.pairplot(mobile_data, vars=numerical_cols[:5], hue="Model", diag_kind="kde", palette="husl")

# Adjust title position
pairplot_fig.fig.suptitle("📊 Scatterplot Matrix with Model as Hue", fontsize=14, fontweight="bold", color="darkblue",
plt.show()
```





- IQOO Neo 7 5G
- vivo Y56 5G
- Infinix Note 12 5G
- Infinix HOT 12 Play
- vivo Y75
- realme Narzo 50A Prime
- realme NARZO 50A PRIME
- Google Pixel 7
- POCO F1
- vivo Y12G
- SAMSUNG Galaxy A23
- realme Narzo 50A
- vivo Y100 5G
- Redmi 9A Sport
- REDMI Note 11
- vivo V25 5G
- SAMSUNG Galaxy S21 FE 5G
- REDMI Note 11T 5G
- realme X3 SuperZoom
- OPPO F21 Pro
- REDMI Note 11S
- OnePlus 8
- vivo Y1s
- OnePlus Nord
- Nokia C01 Plus
- REDMI 10A SPORT
- SAMSUNG Galaxy A54 5G
- vivo Y33s
- REDMI Note 10S
- APPLE iPhone 11
- vivo T1 Pro 5G
- Infinix Smart 5A
- MOTOROLA e22s
- Infinix Hot 12 Pro
- REDMI 9i
- REDMI Note 9
- vivo Y565G
- SAMSUNG M53 5G
- REDMI Note 10 Lite
- APPLE iPhone 14 Plus
- vivo V25 Pro 5G
- Infinix Hot 11
- Infinix Note 12 Pro
- Tecno Pova 3
- REDMI 12c
- LAVA Z2
- Infinix Note 11s Free Fire Edition
- realme C11 2021
- vivo Y21T
- realme 9 5G
- SAMSUNG Galaxy A04
- OPPO F21s Pro
- APPLE iPhone 12
- SAMSUNG Galaxy A13
- Nokia G11 Plus
- SAMSUNG Galaxy A04e
- SAMSUNG Galaxy S23 5G
- LAVA Z21
- Tecno Spark 8T
- OPPO A77

Relationship Between Price and Features (Violin Plot)

```
In [28]: plt.figure(figsize=(16, 6)) # Increase figure size


# Create the violin plot
sns.violinplot(x="Model", y="Prize", data=mobile_data, palette="viridis")

# Improve x-axis readability
plt.xticks(rotation=90, ha="right", fontsize=8) # Rotate labels and adjust font size
plt.xlabel("Model", fontsize=12, fontweight="bold") # X-axis label
plt.ylabel("Prize", fontsize=12, fontweight="bold") # Y-axis label

# Title formatting
plt.title("💰 Price Distribution by Model", fontsize=14, fontweight="bold", color="darkgreen", pad=20)

# Show the plot
plt.show()
```

Price Distribution by Model



Pairwise Feature Dependencies (FacetGrid)

```
In [ ]: g = sns.FacetGrid(mobile_data, col="Model", col_wrap=4, height=4, aspect=1)
g.map_dataframe(sns.scatterplot, x="Battery_", y="Prize", hue="RAM", palette="corm")
g.add_legend()
plt.subplots_adjust(top=0.9)
g.fig.suptitle("🔋 Battery vs. Price Colored by RAM", fontsize=14, fontweight="bold", color="darkblue")
plt.show()
```

```

-----
KeyError                                Traceback (most recent call last)
File ~\anaconda3\Lib\site-packages\seaborn\palettes.py:235, in color_palette(palette, n_colors, desat, as_cmap)
    233 try:
    234     # Perhaps a named matplotlib colormap?
--> 235     palette = mpl_palette(palette, n_colors, as_cmap=as_cmap)
    236 except (ValueError, KeyError): # Error class changed in mpl36

File ~\anaconda3\Lib\site-packages\seaborn\palettes.py:406, in mpl_palette(name, n_colors, as_cmap)
    405 else:
--> 406     cmap = get_colormap(name)
    408 if name in MPL_QUAL_PALS:

File ~\anaconda3\Lib\site-packages\seaborn\_compat.py:62, in get_colormap(name)
    61 try:
--> 62     return mpl.colormaps[name]
    63 except AttributeError:

File ~\anaconda3\Lib\site-packages\matplotlib\cm.py:93, in ColormapRegistry.__getitem__(self, item)
    92 except KeyError:
--> 93     raise KeyError(f"{item!r} is not a known colormap name") from None

```

KeyError: "'corm' is not a known colormap name"

During handling of the above exception, another exception occurred:

```

ValueError                                Traceback (most recent call last)
Cell In[56], line 2
      1 g = sns.FacetGrid(mobile_data, col="Model", col_wrap=4, height=4, aspect=1)
----> 2 g.map_dataframe(sns.scatterplot, x="Battery_", y="Prize", hue="RAM", palette="corm")
      3 g.add_legend()
      4 plt.subplots_adjust(top=0.9)

File ~\anaconda3\Lib\site-packages\seaborn\axisgrid.py:825, in FacetGrid.map_dataframe(self, func, *args, **kwargs)
    822 kwargs["data"] = data_ijk
    824 # Draw the plot
--> 825 self._facet_plot(func, ax, args, kwargs)
    827 # For axis labels, prefer to use positional args for backcompat
    828 # but also extract the x/y kwargs and use if no corresponding arg
    829 axis_labels = [kwargs.get("x", None), kwargs.get("y", None)]

File ~\anaconda3\Lib\site-packages\seaborn\axisgrid.py:854, in FacetGrid._facet_plot(self, func, ax, plot_args, plot_kwa

```

```

rgs)
    852     plot_args = []
    853     plot_kwargs["ax"] = ax
--> 854 func(*plot_args, **plot_kwargs)
    855 # Sort out the supporting information
    856 self._update_legend_data(ax)

```

File ~\anaconda3\Lib\site-packages\seaborn\relational.py:621, in scatterplot(data, x, y, hue, size, style, palette, hue_order, hue_norm, sizes, size_order, size_norm, markers, style_order, legend, ax, **kwargs)

```

    606 def scatterplot(
    607     data=None, *,
    608     x=None, y=None, hue=None, size=None, style=None,
    (... )
    612     **kwargs
    613 ):
    614     p = _ScatterPlotter(
    615         data=data,
    616         variables=dict(x=x, y=y, hue=hue, size=size, style=style),
    617         legend=legend
    618     )
--> 621     p.map_hue(palette=palette, order=hue_order, norm=hue_norm)
    622     p.map_size(sizes=sizes, order=size_order, norm=size_norm)
    623     p.map_style(markers=markers, order=style_order)

```

File ~\anaconda3\Lib\site-packages\seaborn_base.py:838, in VectorPlotter.map_hue(self, palette, order, norm, saturation)

```

    837 def map_hue(self, palette=None, order=None, norm=None, saturation=1):
--> 838     mapping = HueMapping(self, palette, order, norm, saturation)
    839     self._hue_map = mapping

```

File ~\anaconda3\Lib\site-packages\seaborn_base.py:141, in HueMapping.__init__(self, plotter, palette, order, norm, saturation)

```

    138 if map_type == "numeric":
    140     data = pd.to_numeric(data)
--> 141     levels, lookup_table, norm, cmap = self.numeric_mapping(
    142         data, palette, norm,
    143     )
    144 # --- Option 2: categorical mapping using seaborn palette
    147 elif map_type == "categorical":

```

File ~\anaconda3\Lib\site-packages\seaborn_base.py:279, in HueMapping.numeric_mapping(self, data, palette, norm)

```

    277     cmap = palette

```



```

278 else:
--> 279     cmap = color_palette(palette, as_cmap=True)
281 # Now sort out the data normalization
282 if norm is None:

File ~\anaconda3\Lib\site-packages\seaborn\palettes.py:237, in color_palette(palette, n_colors, desat, as_cmap)
235     palette = mpl_palette(palette, n_colors, as_cmap=as_cmap)
236     except (ValueError, KeyError): # Error class changed in mpl36
--> 237         raise ValueError(f"{palette!r} is not a valid palette name")
239 if desat is not None:
240     palette = [desaturate(c, desat) for c in palette]

ValueError: 'corm' is not a valid palette name

```

3D Scatter Plot for Multivariate Dependencies

```

In [26]: # Remove commas and convert to numeric
mobile_data["Prize"] = mobile_data["Prize"].astype(str).str.replace(",", "").astype(float)

# Now, create the scatter plot
fig = plt.figure(figsize=(10, 7))
ax = fig.add_subplot(111, projection="3d")


# Scatter plot with corrected data
sc = ax.scatter(
    mobile_data["Battery_"],
    mobile_data["RAM"],
    mobile_data["Prize"],
    c=mobile_data["Prize"],
    cmap="plasma",
    alpha=0.8
)

# Labels
ax.set_xlabel("Battery Capacity (mAh)")
ax.set_ylabel("RAM (GB)")
ax.set_zlabel("Price (INR)")

# Colorbar
plt.colorbar(sc, label="Price (INR)")

```

```
# Show plot  
plt.show()
```



```
In [57]: # Select only numeric columns  
numeric_cols = mobile_data.select_dtypes(include=['number'])  
  
# Calculate IQR only for numeric columns
```

```

Q1 = numeric_cols.quantile(0.25)
Q3 = numeric_cols.quantile(0.75)
IQR = Q3 - Q1

# Define outlier removal condition
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Remove outliers
mobile_data_cleaned = mobile_data[
    ~((numeric_cols < lower_bound) | (numeric_cols > upper_bound)).any(axis=1)
]

# Display the cleaned dataset
print(mobile_data_cleaned.shape)

# Convert `Prize` to Numeric
mobile_data["Prize"] = mobile_data["Prize"].astype(str).str.replace(",", "").astype(float)

# Rename the Rear Camera and Front Camera column
mobile_data.rename(columns={"Rear Camera": "Rear Camera [MP]"}, inplace=True)
mobile_data.rename(columns={"Front Camera": "Front Camera [MP]"}, inplace=True)

# Display updated column names
print(mobile_data.columns)

# 4. Convert Camera Columns to Numeric
mobile_data["Rear Camera [MP]"] = mobile_data["Rear Camera [MP]"].str.replace("MP", "").astype(float)
mobile_data["Front Camera [MP]"] = mobile_data["Front Camera [MP]"].str.replace("MP", "").astype(float)

# 6. Normalize Numeric Features
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
numeric_cols = ["Memory", "RAM", "Battery_", "Rear Camera [MP]", "Front Camera [MP]", "Mobile Height", "Prize"]
mobile_data[numeric_cols] = scaler.fit_transform(mobile_data[numeric_cols])

# Display the cleaned dataset
mobile_data.head()

```

```
(332, 11)
Index(['Model', 'Colour', 'Memory', 'RAM', 'Battery_', 'Rear Camera [MP]',
      'Front Camera [MP]', 'AI Lens', 'Mobile Height', 'Processor_', 'Prize'],
      dtype='object')
```

Out[57]:

	Model	Colour	Memory	RAM	Battery_	Rear Camera [MP]	Front Camera [MP]	AI Lens	Mobile Height	Processor_	Prize
0	Infinix SMART 7	Night Black	0.200000	0.333333	0.838710	0.065	0.083333	1	0.327457	Unisoc Spreadtrum SC9863A1	0.079659
1	Infinix SMART 7	Azure Blue	0.200000	0.333333	0.838710	0.065	0.083333	1	0.327457	Unisoc Spreadtrum SC9863A1	0.079659
2	MOTOROLA G32	Mineral Gray	0.466667	1.000000	0.677419	0.250	0.266667	0	0.324252	Qualcomm Snapdragon 680	0.138351
3	POCO C50	Royal Blue	0.066667	0.000000	0.677419	0.040	0.083333	0	0.322115	Mediatek Helio A22	0.059054
4	Infinix HOT 30i	Marigold	0.466667	1.000000	0.677419	0.250	0.083333	1	0.327457	G37	0.100888

```
In [59]: from sklearn.preprocessing import LabelEncoder

# List of categorical columns
categorical_cols = ["Model", "Colour", "Processor_"]

# Apply Label Encoding to each categorical column
label_encoders = {}

for col in categorical_cols:
    le = LabelEncoder()
    mobile_data[col] = le.fit_transform(mobile_data[col])
    label_encoders[col] = le # Save encoders for later decoding if needed

# Display the first few rows
print(mobile_data.head())
```


	Model	Colour	Memory	RAM	Battery_	Rear Camera [MP]	\
0	23	159	0.200000	0.333333	0.838710	0.065	
1	23	20	0.200000	0.333333	0.838710	0.065	
2	37	149	0.466667	1.000000	0.677419	0.250	
3	69	201	0.066667	0.000000	0.677419	0.040	
4	12	130	0.466667	1.000000	0.677419	0.250	

	Front Camera [MP]	AI Lens	Mobile Height	Processor_	Prize
0	0.083333	1	0.327457	113	0.079659
1	0.083333	1	0.327457	113	0.079659
2	0.266667	0	0.324252	75	0.138351
3	0.083333	0	0.322115	56	0.059054
4	0.083333	1	0.327457	14	0.100888

```
In [32]: # Compute correlation matrix
correlation_matrix = mobile_data.corr()

# Plot correlation heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Feature Correlation Heatmap")
plt.show()

# Find top correlated features
corr_with_price = correlation_matrix["Prize"].abs().sort_values(ascending=False)
print("Top Correlated Features with Price:\n", corr_with_price)
```



Top Correlated Features with Price:

Prize	1.000000
Memory	0.563535
RAM	0.529474
Front Camera [MP]	0.529013
Rear Camera [MP]	0.406784
Mobile Height	0.168303
AI Lens	0.156336
Model	0.073833
Processor_	0.049600
Battery_	0.046250
Colour	0.040595

Name: Prize, dtype: float64

```
In [61]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split

# Define X (features) and y (target variable)
X = mobile_data.drop(columns=["Prize"]) # Remove target column
y = mobile_data["Prize"]

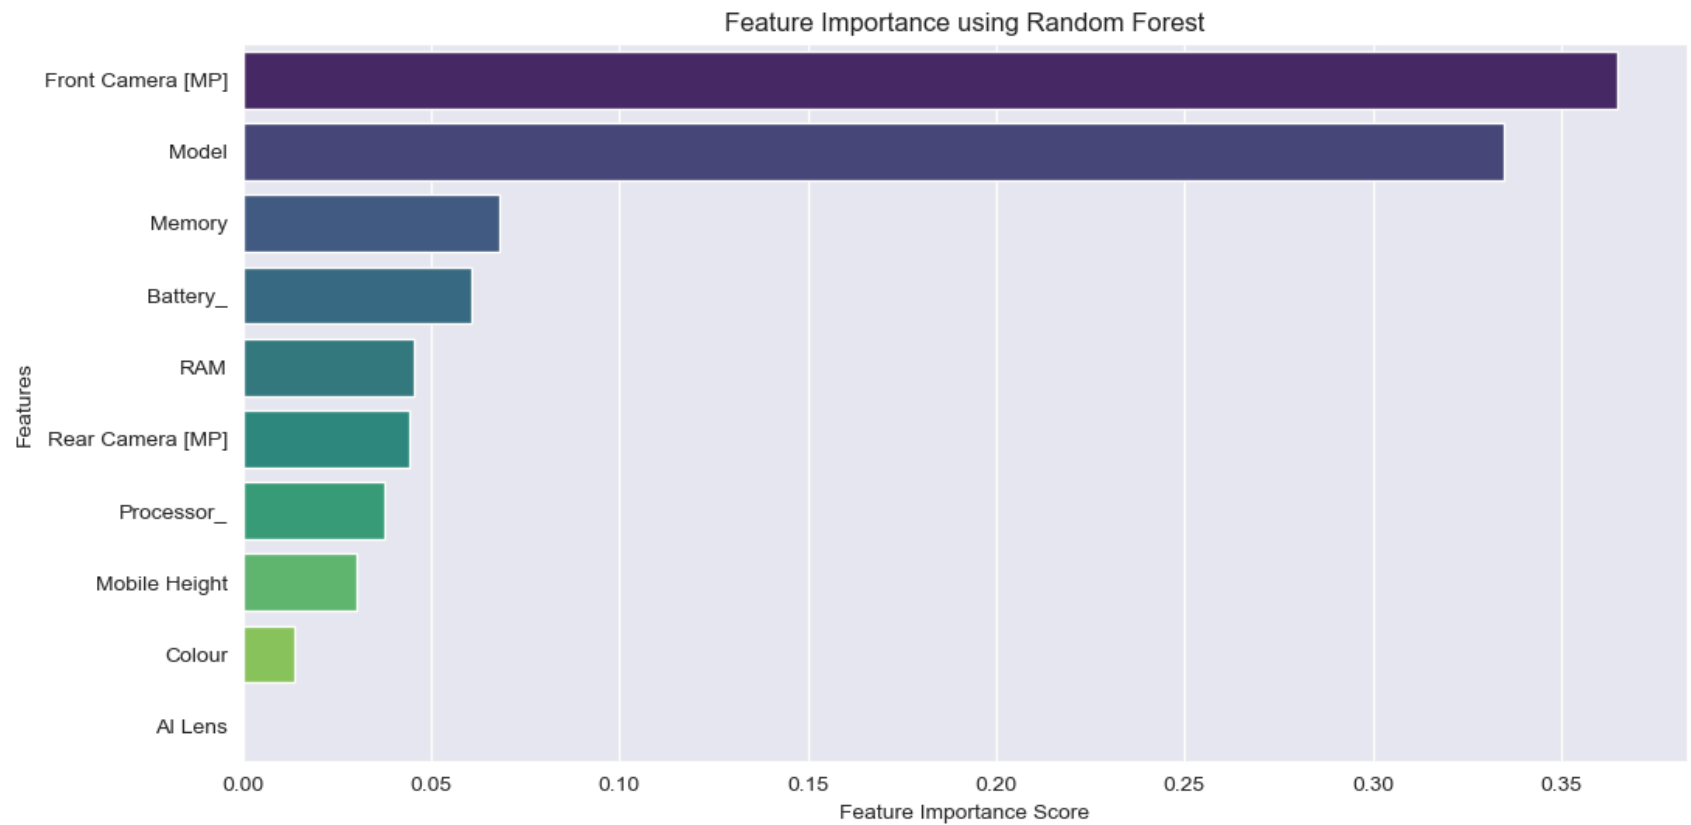
# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train a Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Get feature importances
feature_importances = pd.Series(rf_model.feature_importances_, index=X.columns)
feature_importances = feature_importances.sort_values(ascending=False)

# Plot feature importance
plt.figure(figsize=(12, 6))
sns.barplot(x=feature_importances, y=feature_importances.index, palette="viridis")
plt.xlabel("Feature Importance Score")
plt.ylabel("Features")
plt.title("Feature Importance using Random Forest")
plt.show()

# Print top important features
print("Top Important Features:\n", feature_importances)
```



Top Important Features:

Front Camera [MP]	0.365029
Model	0.334752
Memory	0.068347
Battery_	0.060567
RAM	0.045294
Rear Camera [MP]	0.044153
Processor_	0.037650
Mobile Height	0.030405
Colour	0.013591
AI Lens	0.000211

dtype: float64

```
In [63]: from sklearn.decomposition import PCA
         from sklearn.preprocessing import StandardScaler


         # Standardize the features
```



```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply PCA
pca = PCA(n_components=10) # Reduce to 10 components
X_pca = pca.fit_transform(X_scaled)

# Explained variance ratio
plt.figure(figsize=(10, 5))
plt.plot(range(1, 11), np.cumsum(pca.explained_variance_ratio_), marker="o", linestyle="--")
plt.xlabel("Number of Components")
plt.ylabel("Cumulative Explained Variance")
plt.title("PCA - Explained Variance")
plt.show()
```



```
In [65]: selected_features = feature_importances[:25].index # Keep top 25 features
X_selected = mobile_data[selected_features]
print("Final Selected Features:\n", X_selected.columns)
```

Final Selected Features:

```
Index(['Front Camera [MP]', 'Model', 'Memory', 'Battery_', 'RAM',
      'Rear Camera [MP]', 'Processor_', 'Mobile Height', 'Colour', 'AI Lens'],
      dtype='object')
```

```
In [54]: # Define target variable (Price)
y = mobile_data["Prize"]

# Use only the selected features from the previous step
selected_features = ["RAM", "Memory", "Battery_", "Rear Camera [MP]", "Front Camera [MP]",
```

```

        "Mobile Height", "AI Lens", "Processor_", "Model", "Colour" ] # Example selected features
X = mobile_data[selected_features]

# Train-test split (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Training Set:", X_train.shape)
print("Testing Set:", X_test.shape)

```

Training Set: (424, 10)

Testing Set: (107, 10)

```

In [67]: # Train Linear Regression Model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)

# Predictions
y_pred_lr = lr_model.predict(X_test)

```

```

In [69]: # Train Decision Tree Model
dt_model = DecisionTreeRegressor(random_state=42)
dt_model.fit(X_train, y_train)

# Predictions
y_pred_dt = dt_model.predict(X_test)

```

```

In [71]: # Train Random Forest Model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Predictions
y_pred_rf = rf_model.predict(X_test)

```

```

In [73]: # 📌 Evaluation Function
def evaluate_model(y_test, y_pred, model_name):
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse) # Root Mean Squared Error
    r2 = r2_score(y_test, y_pred)

    print(f"📊 {model_name} Performance:")

```

```

print(f"    - Mean Absolute Error (MAE): {mae:.2f}")
print(f"    - Mean Squared Error (MSE): {mse:.2f}")
print(f"    - Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"    - R2 Score: {r2:.4f}\n")

# 🚀 Evaluate all three models
evaluate_model(y_test, y_pred_lr, "Linear Regression")
evaluate_model(y_test, y_pred_dt, "Decision Tree")
evaluate_model(y_test, y_pred_rf, "Random Forest")

# 🚀 Show actual vs predicted prices (for first 10 values)
predictions_df = pd.DataFrame({
    "Actual Price": y_test.values,
    "Predicted (LR)": y_pred_lr,
    "Predicted (DT)": y_pred_dt,
    "Predicted (RF)": y_pred_rf
})

print("\n🚀 Actual vs Predicted Prices (First 10 values):")
print(predictions_df.head(10))

```



Linear Regression Performance:

- Mean Absolute Error (MAE): 0.06
- Mean Squared Error (MSE): 0.01
- Root Mean Squared Error (RMSE): 0.11
- R² Score: 0.2672



Decision Tree Performance:

- Mean Absolute Error (MAE): 0.03
- Mean Squared Error (MSE): 0.01
- Root Mean Squared Error (RMSE): 0.12
- R² Score: 0.1252



Random Forest Performance:

- Mean Absolute Error (MAE): 0.03
- Mean Squared Error (MSE): 0.01
- Root Mean Squared Error (RMSE): 0.09
- R² Score: 0.5565



Actual vs Predicted Prices (First 10 values):

	Actual Price	Predicted (LR)	Predicted (DT)	Predicted (RF)
0	0.288203	0.233110	0.288203	0.297918
1	0.100888	0.179224	0.100888	0.102635
2	0.537956	0.297063	0.537956	0.543063
3	0.094644	0.092097	0.097766	0.091409
4	0.113376	0.195817	0.113376	0.111153
5	0.113376	0.095531	0.113376	0.114536
6	0.169570	0.209377	0.169570	0.174988
7	0.188302	0.221359	0.144595	0.207124
8	0.213277	0.228175	0.213277	0.243667
9	0.075913	0.048842	0.075913	0.073430

```
In [75]: import matplotlib.pyplot as plt
import seaborn as sns

# Get feature importance from Decision Tree
dt_importance = dt_model.feature_importances_

# Get feature importance from Random Forest
rf_importance = rf_model.feature_importances_

# Convert to DataFrame for visualization
```

```

feature_importance_df = pd.DataFrame({
    "Feature": X.columns,
    "Decision Tree Importance": dt_importance,
    "Random Forest Importance": rf_importance
})

# Sort by importance (Random Forest)
feature_importance_df = feature_importance_df.sort_values(by="Random Forest Importance", ascending=False)

# 📌 Print Top Features
print("📊 Top Features based on Random Forest:")
print(feature_importance_df.head(10))

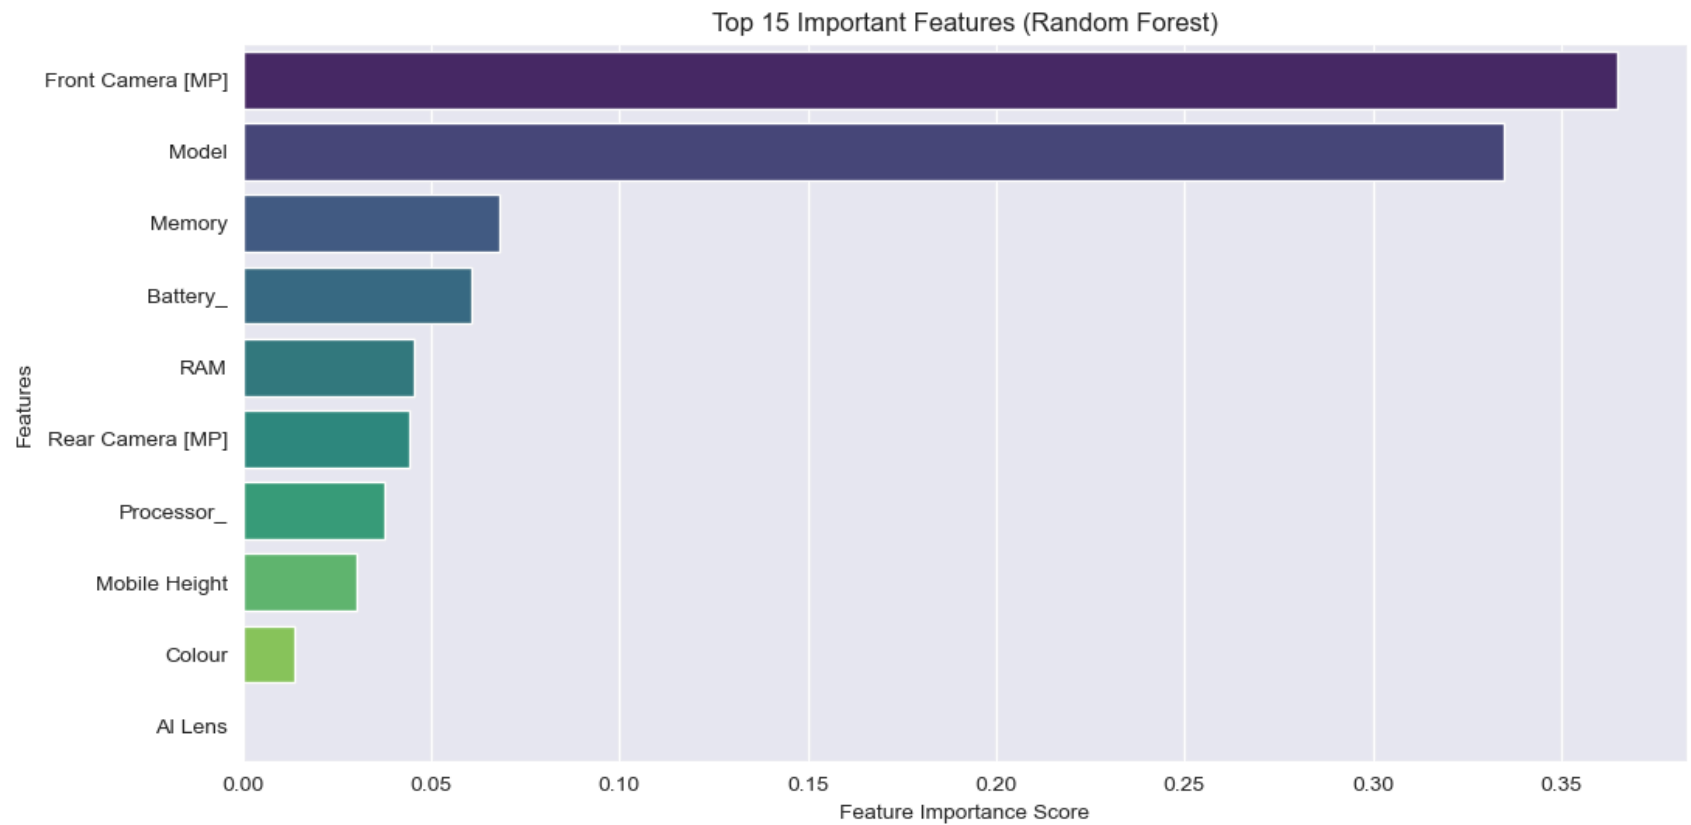
# 📌 Visualization of Feature Importance
plt.figure(figsize=(12, 6))
sns.barplot(x=feature_importance_df["Random Forest Importance"][:15],
            y=feature_importance_df["Feature"][:15],
            palette="viridis")

plt.xlabel("Feature Importance Score")
plt.ylabel("Features")
plt.title("Top 15 Important Features (Random Forest)")
plt.show()

```

📊 Top Features based on Random Forest:

	Feature	Decision Tree Importance	Random Forest Importance
6	Front Camera [MP]	0.323210	0.365029
0	Model	0.385938	0.334752
2	Memory	0.075037	0.068347
4	Battery_	0.043459	0.060567
3	RAM	0.066478	0.045294
5	Rear Camera [MP]	0.032285	0.044153
9	Processor_	0.025824	0.037650
8	Mobile Height	0.022389	0.030405
1	Colour	0.025364	0.013591
7	AI Lens	0.000015	0.000211



```
In [91]: from pptx import Presentation
from pptx.util import Inches

# Create a PowerPoint Presentation
prs = Presentation()

# Add Title Slide
slide_layout = prs.slide_layouts[0]
slide = prs.slides.add_slide(slide_layout)
title = slide.shapes.title
subtitle = slide.placeholders[1]
title.text = "Mobile Phone Price Prediction"
subtitle.text = "By Your Name | Date"

# Add an Introduction Slide
slide_layout = prs.slide_layouts[1]
```

```
slide = prs.slides.add_slide(slide_layout)
title = slide.shapes.title
content = slide.placeholders[1]
title.text = "Introduction"
content.text = "Goal: Predict mobile phone prices using Machine Learning.\nTechniques Used: Data Cleaning, Feature Sele

# Save the PowerPoint
prs.save("Mobile_Price_Prediction_Report.pptx")
print("✅ PowerPoint Created Successfully!")
```

✅ PowerPoint Created Successfully!

In [83]: !pip install python-pptx

Collecting python-pptx

Downloading python_pptx-1.0.2-py3-none-any.whl.metadata (2.5 kB)

Requirement already satisfied: Pillow>=3.3.2 in c:\users\user\anaconda3\lib\site-packages (from python-pptx) (10.3.0)

Collecting XlsxWriter>=0.5.7 (from python-pptx)

Downloading XlsxWriter-3.2.2-py3-none-any.whl.metadata (2.8 kB)

Requirement already satisfied: lxml>=3.1.0 in c:\users\user\anaconda3\lib\site-packages (from python-pptx) (5.2.1)

Requirement already satisfied: typing-extensions>=4.9.0 in c:\users\user\anaconda3\lib\site-packages (from python-pptx) (4.11.0)

Downloading python_pptx-1.0.2-py3-none-any.whl (472 kB)

```
----- 0.0/472.8 kB ? eta -:-:--
----- 0.0/472.8 kB ? eta -:-:--
----- 10.2/472.8 kB ? eta -:-:--
-- ----- 30.7/472.8 kB 325.1 kB/s eta 0:00:02
--- ----- 41.0/472.8 kB 279.3 kB/s eta 0:00:02
----- 61.4/472.8 kB 363.1 kB/s eta 0:00:02
----- 143.4/472.8 kB 607.9 kB/s eta 0:00:01
----- 143.4/472.8 kB 607.9 kB/s eta 0:00:01
----- 225.3/472.8 kB 655.6 kB/s eta 0:00:01
----- 225.3/472.8 kB 655.6 kB/s eta 0:00:01
----- 307.2/472.8 kB 731.4 kB/s eta 0:00:01
----- 307.2/472.8 kB 731.4 kB/s eta 0:00:01
----- 337.9/472.8 kB 655.4 kB/s eta 0:00:01
----- 389.1/472.8 kB 713.5 kB/s eta 0:00:01
----- 399.4/472.8 kB 637.9 kB/s eta 0:00:01
----- 419.8/472.8 kB 624.4 kB/s eta 0:00:01
----- 450.6/472.8 kB 655.2 kB/s eta 0:00:01
----- 450.6/472.8 kB 655.2 kB/s eta 0:00:01
----- 472.8/472.8 kB 603.8 kB/s eta 0:00:00
```

Downloading XlsxWriter-3.2.2-py3-none-any.whl (165 kB)

```
----- 0.0/165.1 kB ? eta -:-:--
-- ----- 10.2/165.1 kB ? eta -:-:--
----- 41.0/165.1 kB 393.8 kB/s eta 0:00:01
----- 163.8/165.1 kB 1.2 MB/s eta 0:00:01
----- 165.1/165.1 kB 1.1 MB/s eta 0:00:00
```

Installing collected packages: XlsxWriter, python-pptx

Successfully installed XlsxWriter-3.2.2 python-pptx-1.0.2

In [93]: `import joblib`

```
# Save the trained Random Forest model
joblib.dump(rf_model, "random_forest_mobile_price.pkl")
```

```
print("Model saved successfully!")
```

Model saved successfully!

```
In [95]: # Load the saved model
rf_model_loaded = joblib.load("random_forest_mobile_price.pkl")

print("Model loaded successfully!")
```

Model loaded successfully!

```
In [111]: # Ensure new_data has the same features as the training data
new_data.columns = X_train.columns
print(new_data.columns)
new_data = pd.DataFrame([[23, 20, 0.200000, 0.333333, 0.838710, 0.065, 0.083333, 1, 0.327457, 113]]) # Example data

# Predict the price using the loaded model
predicted_price = rf_model_loaded.predict(new_data)

print("Predicted Mobile Price:", predicted_price[0])
```

```
Index(['Model', 'Colour', 'Memory', 'RAM', 'Battery_', 'Rear Camera [MP]',
      'Front Camera [MP]', 'AI Lens', 'Mobile Height', 'Processor_'],
      dtype='object')
```

Predicted Mobile Price: 0.08264601206308753

```
In [99]: mobile_data
```

Out[99]:

	Model	Colour	Memory	RAM	Battery_	Rear Camera [MP]	Front Camera [MP]	AI Lens	Mobile Height	Processor_	Prize
0	23	159	0.200000	0.333333	0.838710	0.065	0.083333	1	0.327457	113	0.079659
1	23	20	0.200000	0.333333	0.838710	0.065	0.083333	1	0.327457	113	0.079659
2	37	149	0.466667	1.000000	0.677419	0.250	0.266667	0	0.324252	75	0.138351
3	69	201	0.066667	0.000000	0.677419	0.040	0.083333	0	0.322115	56	0.059054
4	12	130	0.466667	1.000000	0.677419	0.250	0.083333	1	0.327457	14	0.100888
...
536	118	49	1.000000	1.000000	0.500000	0.250	0.200000	0	0.293536	89	0.987512
537	32	52	0.066667	0.000000	0.370968	0.025	0.033333	0	0.219017	68	0.063412
538	123	259	0.200000	0.333333	0.677419	0.250	0.133333	0	0.327457	35	0.113263
539	110	17	0.466667	1.000000	0.677419	0.250	0.533333	0	0.314103	11	0.475518
540	59	215	0.466667	0.333333	0.677419	0.250	0.133333	0	0.324786	57	0.188302

531 rows × 11 columns

```
In [113]: actual_price = mobile_data["Prize"][0] # Assuming this is the actual price
predicted_price = predicted_price[0]

error = actual_price - predicted_price
print(f"Actual Price: {actual_price}")
print(f"Predicted Price: {predicted_price}")
print(f"Prediction Error: {error}")
```

Actual Price: 0.07965883689856267

Predicted Price: 0.08264601206308753

Prediction Error: -0.0029871751645248606

```
In [105]: print(X_train.columns)
```

```
Index(['Model', 'Colour', 'Memory', 'RAM', 'Battery_', 'Rear Camera [MP]',  
      'Front Camera [MP]', 'AI Lens', 'Mobile Height', 'Processor_'],  
      dtype='object')
```

```
In [117]: # Import necessary Libraries  
from sklearn.ensemble import GradientBoostingRegressor  
from sklearn.metrics import mean_absolute_error, mean_squared_error  
  
# Initialize the Gradient Boosting model  
gb_model = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)  
  
# Train the model  
gb_model.fit(X_train, y_train)  
  
# Predict on test data  
y_pred_gb = gb_model.predict(X_test)  
  
# Print first 10 predicted vs actual values  
print("\n📊 Actual vs Predicted Prices (Gradient Boosting - First 10 values):")  
print(pd.DataFrame({  
    "Actual Price": y_test[:10].values,  
    "Predicted (GB)": y_pred_gb[:10]  
}))  
  
# Evaluate the model  
mae = mean_absolute_error(y_test, y_pred_gb)  
rmse = mean_squared_error(y_test, y_pred_gb, squared=False)  
  
print("\n📊 Model Evaluation (Gradient Boosting Regressor):")  
print(f"✅ Mean Absolute Error (MAE): {mae:.4f}")  
print(f"✅ Root Mean Squared Error (RMSE): {rmse:.4f}")
```

Actual vs Predicted Prices (Gradient Boosting - First 10 values):

	Actual Price	Predicted (GB)
0	0.288203	0.286618
1	0.100888	0.111670
2	0.537956	0.537126
3	0.094644	0.079139
4	0.113376	0.117566
5	0.113376	0.105173
6	0.169570	0.166106
7	0.188302	0.213113
8	0.213277	0.254080
9	0.075913	0.074726

Model Evaluation (Gradient Boosting Regressor):

✓ Mean Absolute Error (MAE): 0.0347

✓ Root Mean Squared Error (RMSE): 0.0965

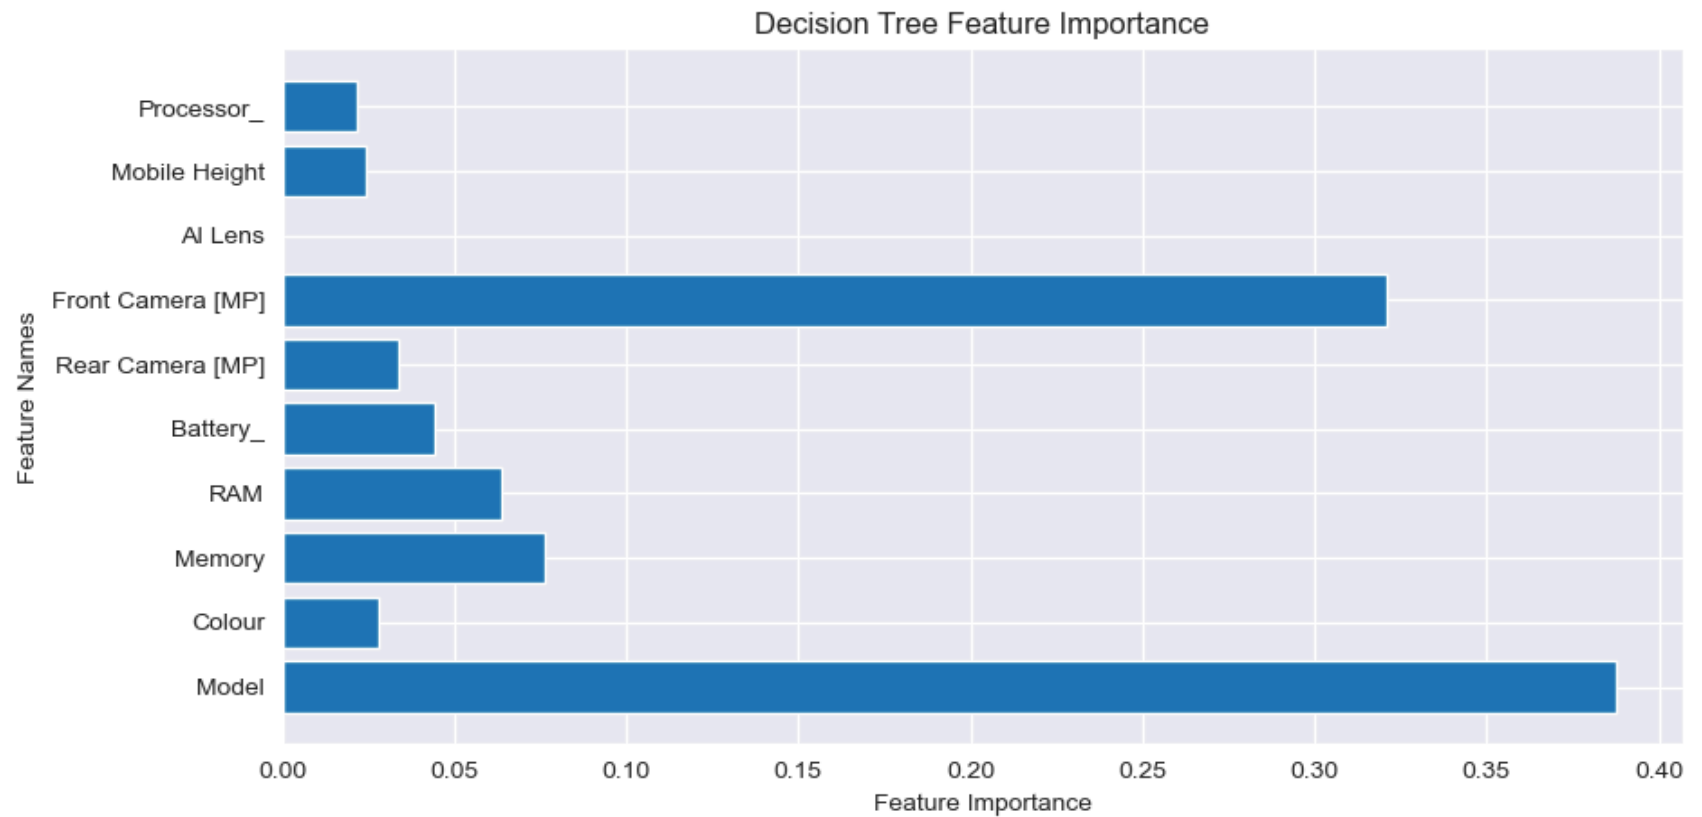
```
In [119]: dt_model = DecisionTreeRegressor(max_depth=10, min_samples_split=5, min_samples_leaf=2, random_state=42)
          dt_model.fit(X_train, y_train)
```

```
Out[119]: DecisionTreeRegressor
          DecisionTreeRegressor(max_depth=10, min_samples_leaf=2, min_samples_split=5,
                                random_state=42)
```

```
In [121]: import joblib
          joblib.dump(dt_model, "final_mobile_price_model.pkl")
```

```
Out[121]: ['final_mobile_price_model.pkl']
```

```
In [123]: import matplotlib.pyplot as plt
          feature_importance = dt_model.feature_importances_
          plt.figure(figsize=(10,5))
          plt.barh(X_train.columns, feature_importance)
          plt.xlabel("Feature Importance")
          plt.ylabel("Feature Names")
          plt.title("Decision Tree Feature Importance")
          plt.show()
```



```
In [125]: loaded_model = joblib.load("final_mobile_price_model.pkl")
predicted_price = loaded_model.predict(new_data)
print("Predicted Mobile Price:", predicted_price[0])
```

Predicted Mobile Price: 0.09132286215188588

```
In [ ]:
```

```
In [146]: import os
os.system("jupyter nbconvert --to pdf mobile_data.ipynb")
```

Out[146]: 1

```
In [141]: pip install pandoc
```

Requirement already satisfied: pandoc in c:\users\user\anaconda3\lib\site-packages (2.4)
Requirement already satisfied: plumbum in c:\users\user\anaconda3\lib\site-packages (from pandoc) (1.9.0)
Requirement already satisfied: ply in c:\users\user\anaconda3\lib\site-packages (from pandoc) (3.11)
Requirement already satisfied: pywin32 in c:\users\user\anaconda3\lib\site-packages (from plumbum->pandoc) (305.1)
Note: you may need to restart the kernel to use updated packages.

In [143]:

Cell In[143], line 1

```
sudo apt install texlive-xetex texlive-fonts-recommended
```

^

SyntaxError: invalid syntax

In []: